

# Max-flow Parallelization in OpenMP vs. DSLs

Sophia Zalewski, Scott Routledge

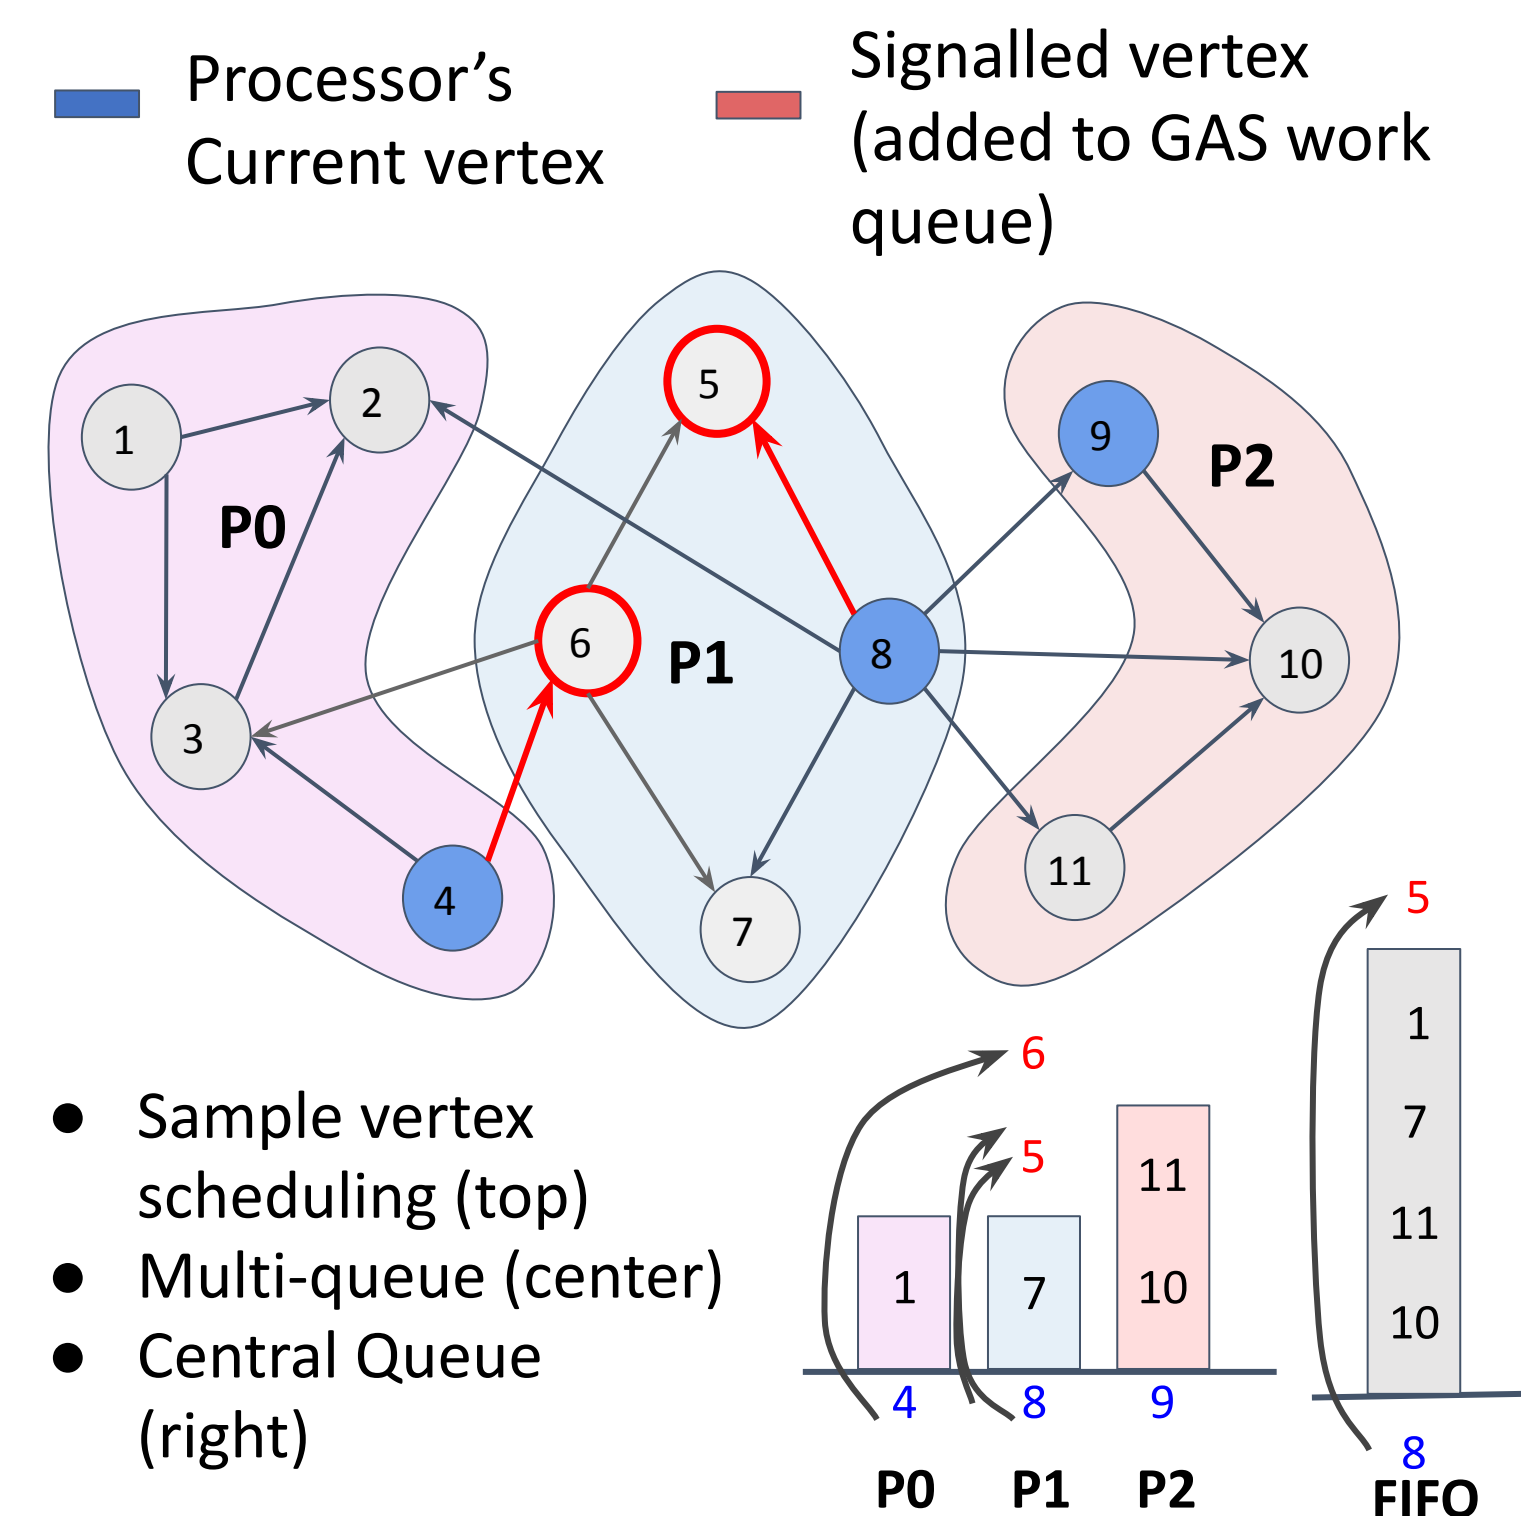
## Abstract

In this project, we compare the advantages and disadvantages of Domain Specific Languages (DSLs) for max-flow problems. We:

- Design our own compact DSL for graph vertex problems (**GraphLabLite**)
- Implement three different algorithms to solve max flow problems (Ford Fulkerson, Dinic's, Push-Relabel), and compare the difference between the OpenMP-parallelized algorithms and DSL-parallelized algorithm

## Work Scheduling

- **Vertex Partitioning** - vertices in the same neighborhood (ideally) assigned to the same processor, each proc does GAS on its assigned vertices, using locks / critical sections when needed depending on consistency model
- **Simultaneous** - GAS on every vertex, every iter
- **Signaling** - current vertex "signals" which neighbor vertices GAS is done on next (if any), user supplies signal call
  - Central Queue, Multi-Queue



## GraphLabLite (Our DSL)

- Iterative vertex-centric operations, runs **Gather** (accumulate data from neighbors), **Apply** (update vertex data) **Scatter** (distribute updated data to neighbors) until converge
- Inspired by CMU's GraphLab implementation

## Graph Representation

**data** stored at each edge and vertex + read-only global data, **locks** for each vertex and edge, **proc info**

```
struct tVertex{  
    • void *data  
    • lock v_lock  
    • Int proc_id  
    • vec<int> border_edges_in,  
      border_edges_out  
};  
  
struct tEdge{  
    • Int u,v  
    • void *data  
    • lock e_lock  
};  
  
class tGraph{  
    • Int num nodes  
    • vec<tEdge> edges  
    • vec<int> in_edges  
    • vec<int> out_edges  
    • void *global_data  
};
```


Which neighboring edges belong to a different processor (for enforcing consistency)

Which **processor** this vertex is assigned to

```
tGraph::gather_context INGOING;  
tGraph::scatter_context OUTGOING;  
tGraph::consist = FULL;  
tGraph::schedule = SIMULTANEOUS;
```

User picks edges to do scatter/gather on, specifies schedule and consistency model


## Consistency Models



- **Vertex** - no locks or critical sections during GAS
- **Edge** - Lock this vertex's "border edges"
- **Full** - Lock border edges + neighbor vertices who are assigned to a different processor or can be accessed by another processor (has non-empty border edges)

## Generating Tests (Method)

- Wrote a **flow-graph generating python script**, generates 5 different graph types (user specifies num nodes and layers + type)

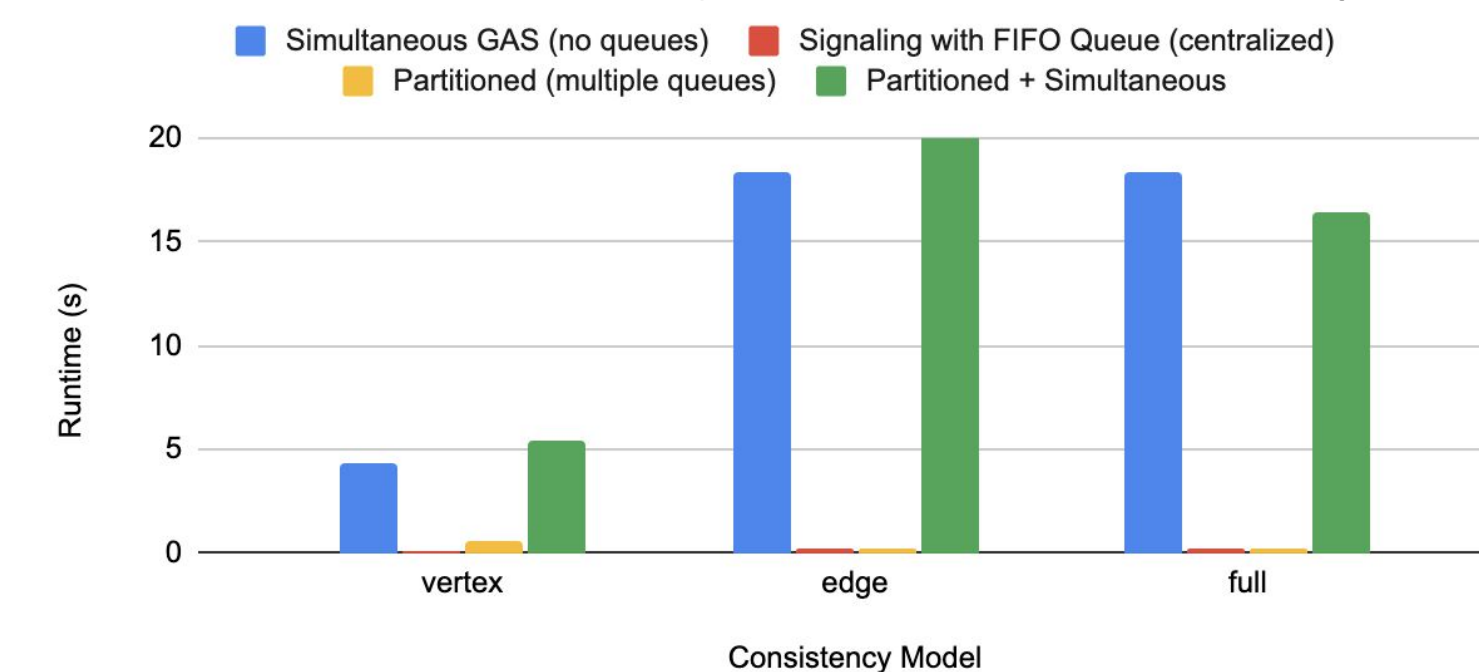


(1) Dense Layer Graph, (2) Sparse Layer Graph (3) Exclusive Path Graph, (4) Line Graph, (5) Unbalanced Graph

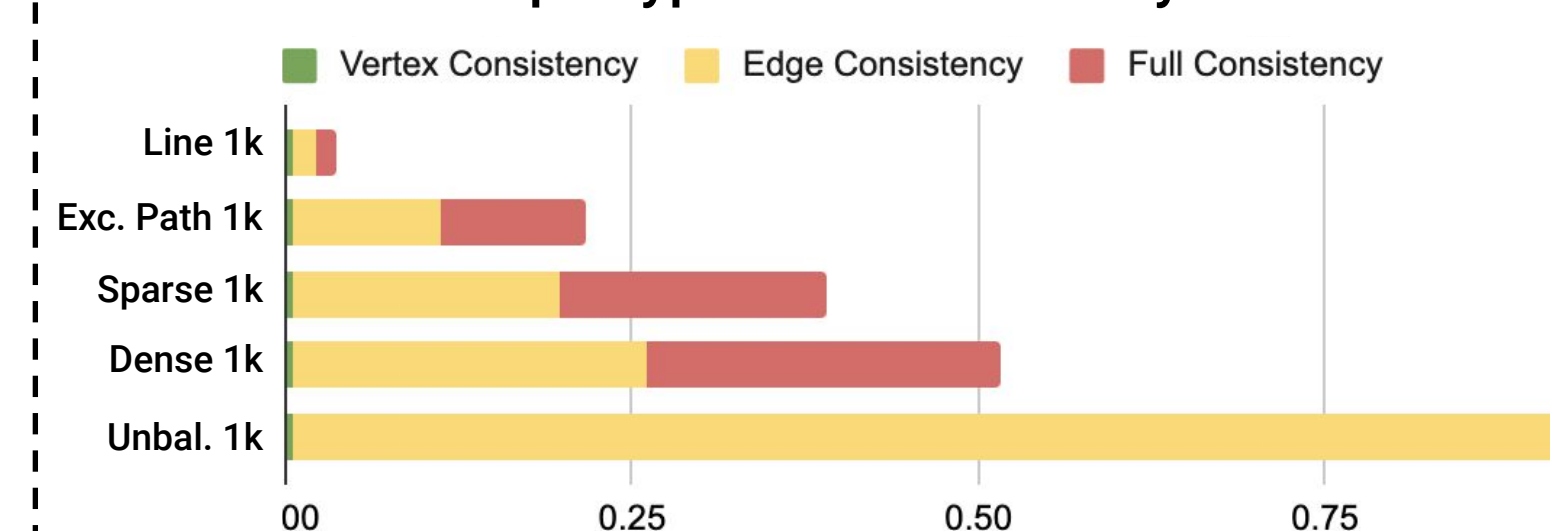
## Push Relabel in GraphLabLite

- Push flow to vertices of lower "height" or relabel / increase height until graph converges


## Runtime vs. Scheduling for all three consistency models



## Vertices Requiring Locks/Critical Sections (%) by Graph Type and Consistency Model



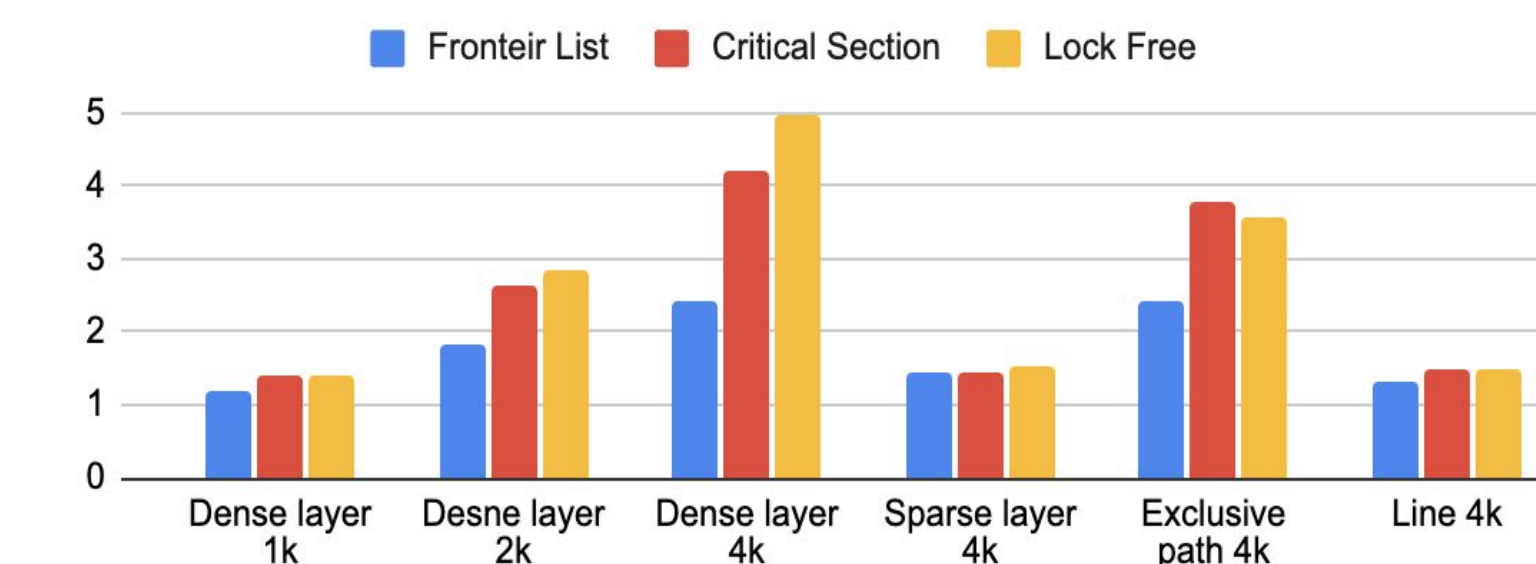
## Vertices Per Processor (%) by Graph Type



## Ford Fulkerson in OpenMP

- Uses BFS to repeatedly push flow from source to sink (parallelize the BFS)
- **Lock-Free** BFS fastest on all graph types
- Fastest speedup on dense layer graph with 4k nodes, 120 levels (**5x speedup**)


## Ford Fulkerson Speedup vs. BFS Implementation



## Dinic's in OpenMP

- BFS to make a "level graph", then DFS to push "blocking flows"
- Dinic's BFS (adjacency list) reaches around **4x speedup** (weak scaling) on 32 threads => **2x Dinic's speedup**

## Dinic's BFS Speedup vs. Implementation, Scaling, #Threads



## Comparison

- Overhead to learning DSL, but easier to program
- Easy to try different schedules w/o worrying about implementation
- Actual GraphLab implementation would run even faster than ours

Impl.	dense-5k runtime (s)	dense 40k-runtime (s)
FF Par	16.08770	—
Dinics Par	0.06737	0.141852
Push Relabel	0.80958	7.80220